

# MuCoSim: Analysis of MiniMD

Jinwen Pan

Advisors: Jan Eitzinger and Dominik Ernst

July 18, 2022

01 Project Description

02 Scalability Comparison of OpenMP and MPI

} Talk I

03 Whole Application Measurements

04 Runtime Profile

05 Hot Spots Measurements

} Talk II

## Molecular Dynamics Simulation Mini-App

MiniMD is a molecular dynamics proxy application simulating a multiparticle system based on Newtonian Mechanics.

- Force Calculation: Lennard-Jones Potential and Embedded Atom Model Potential
- Cutoff Distance: Half/Full Neighbor Lists Approach and Cell Lists Approach

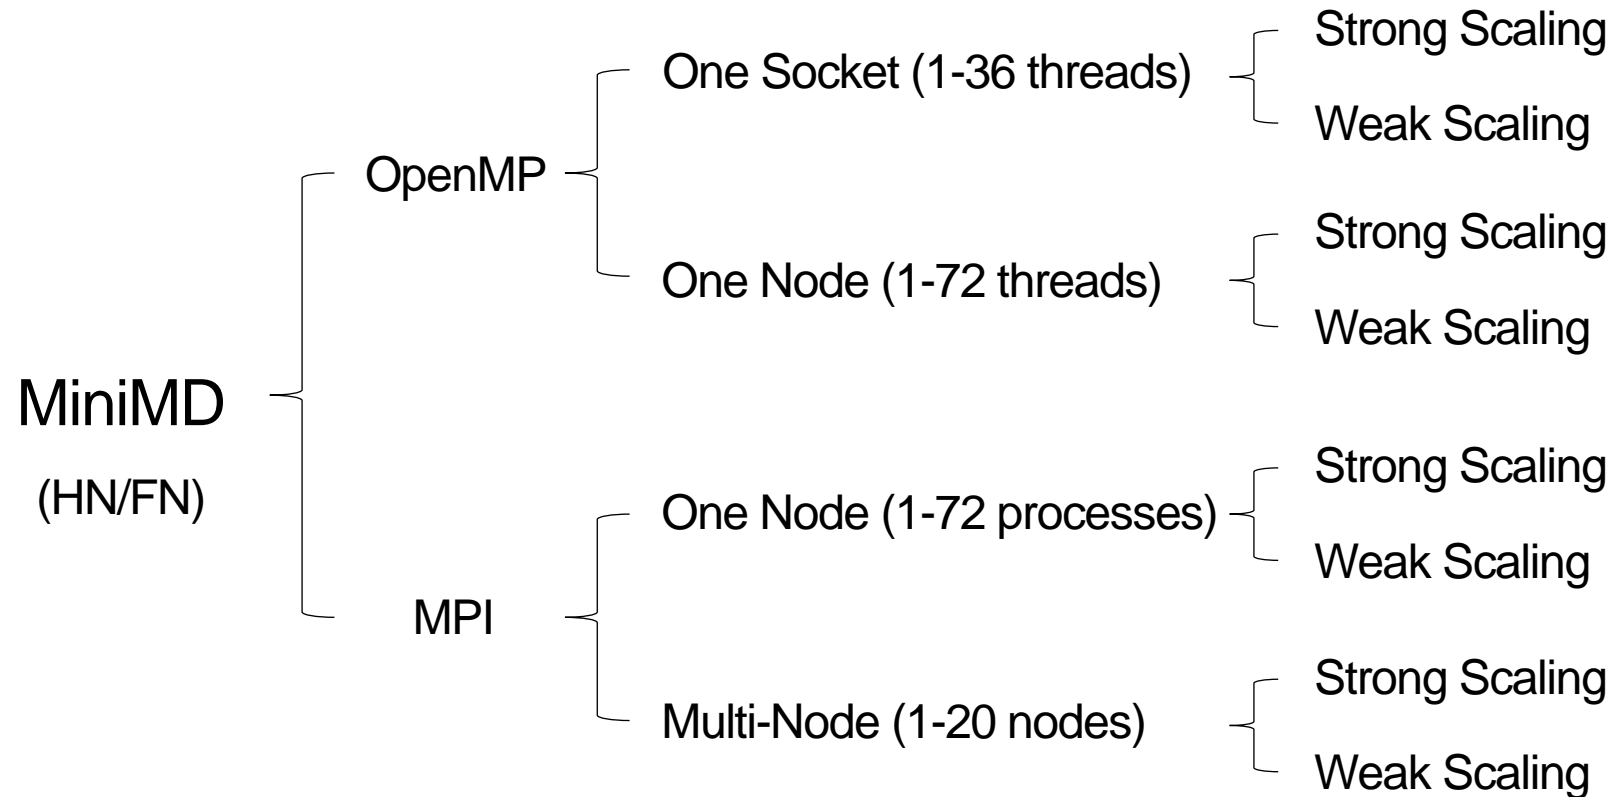
## Testsystem Fritz

- CPU Name: Intel(R) Xeon(R) Platinum 8360Y CPU @ 2.40GHz
- CPU Type: Intel Icelake SP Processor
- 944 compute nodes with 2 sockets per node and 36 cores per socket
- 4 NUMA domains (2+2) per node: 18 cores per domain (Cluster On Die / SNC active)
- One thread per core: SMT is not active
- Each core has an individual L1 cache (48 KB) and L2 cache (1.25 MB)
- Each socket has an individual L3 cache (54 MB)

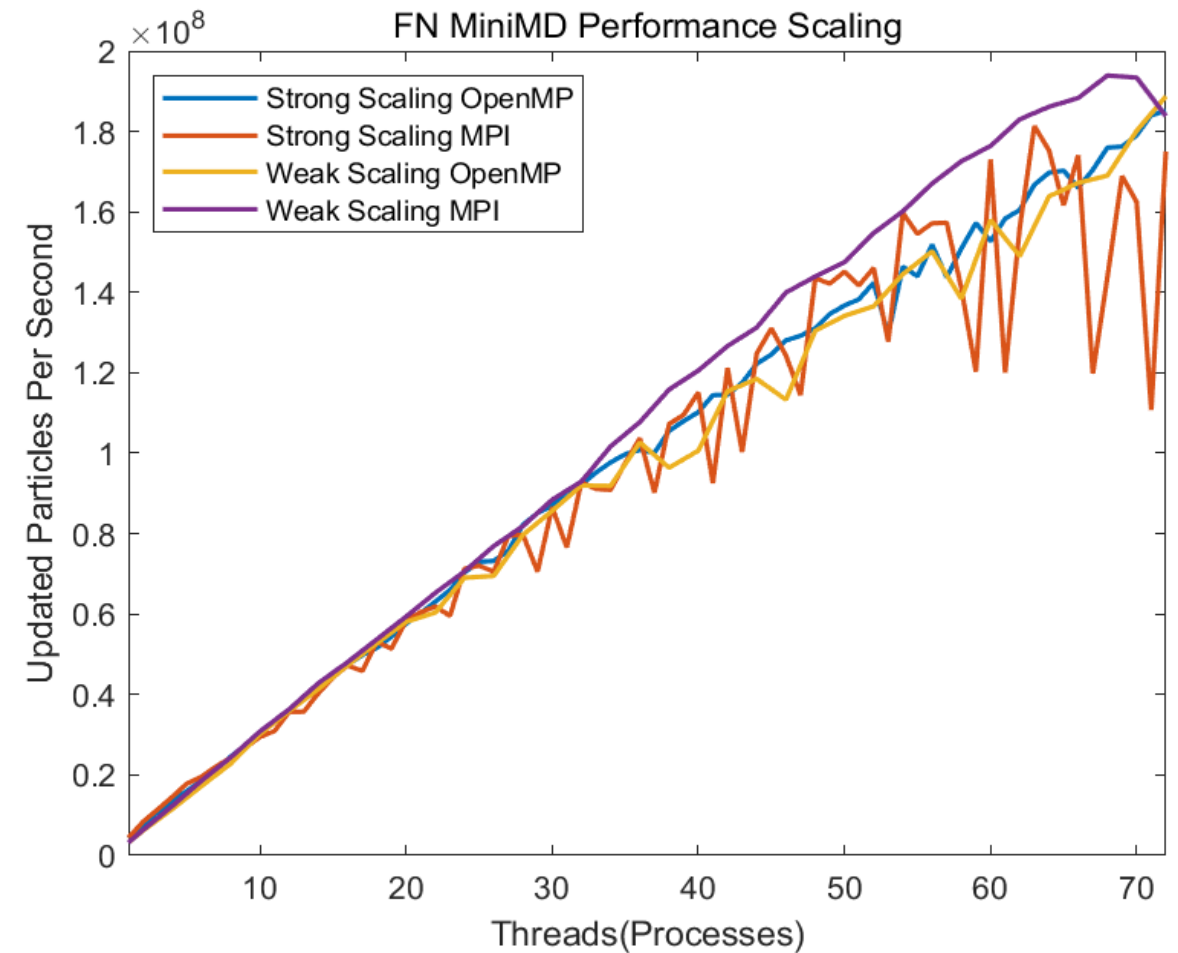
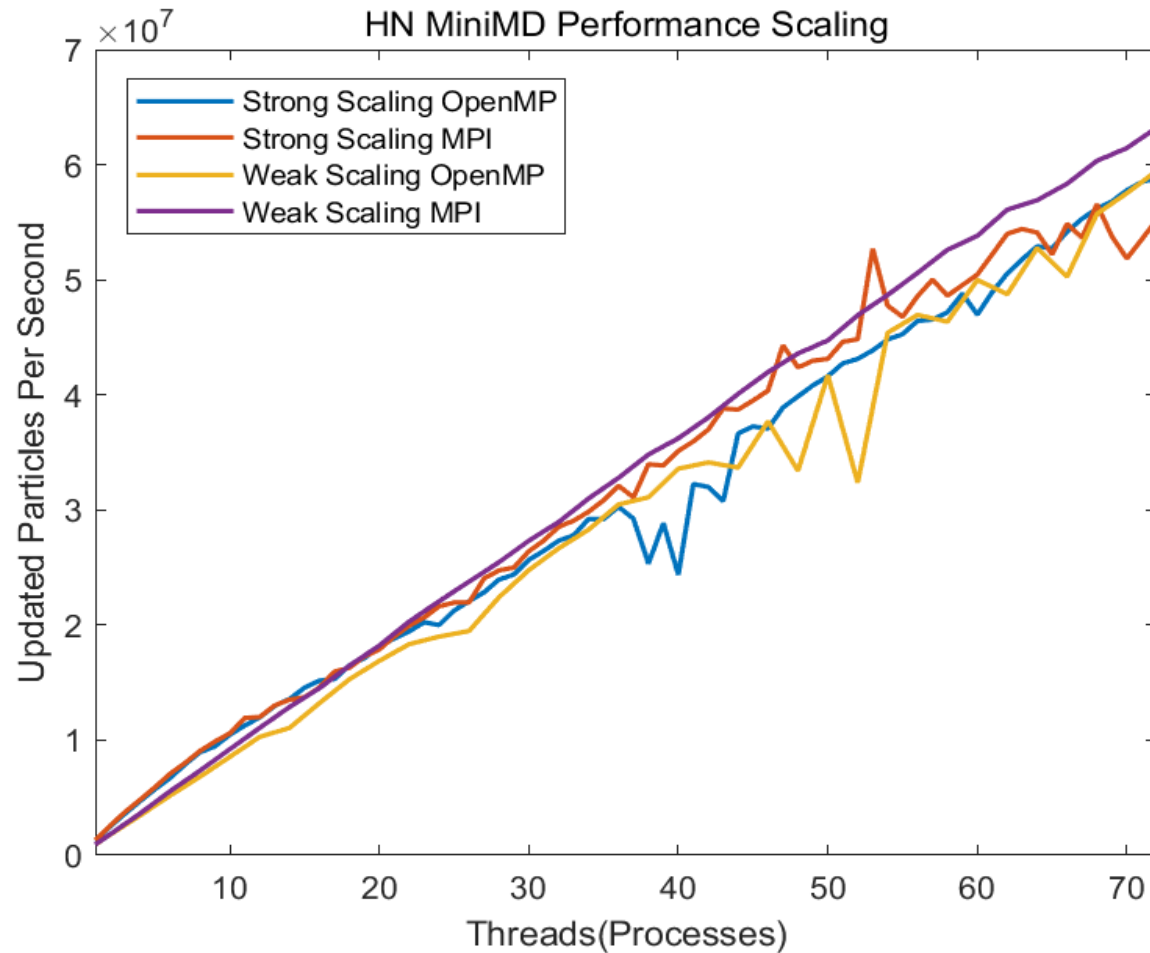
## Modules Loaded

- Intel/2021.4.0
- Intelmpi/2021.4.0
- Likwid/5.2.1

# Scalability Comparison of OpenMP and MPI



# Scalability Comparison of OpenMP and MPI



## Single Instruction Multiple Data (SIMD) Vectorization

SIMD is a feature of a single core, but applies a same instruction to multiple operands in parallel with wide registers.

- Scalar: register width = 64 bits (32 bits)
  - > 1 double (1 single) precision FP operand
- SSE: register width = 128 bits
  - > 2 double (4 single) precision FP operands
- AVX: register width = 256 bits
  - > 4 double (8 single) precision FP operands
- AVX512: register width = 512 bits
  - > 8 double (16 single) precision FP operands

## Compilation Configuration

#OPTS = -Ofast -no-vec

#OPTS = -Ofast -xSSE4.2 -DUSE\_SIMD

#OPTS = -Ofast -xCORE-AVX2 -DUSE\_SIMD

#OPTS = -Ofast -xCORE-AVX512 -qopt-zmm-usage=high -DUSE\_SIMD

## FLOPS\_DP Group

\$ srun --constraint=hwperf likwid-perfctr -g FLOPS\_DP -C SO:1 ../miniMD-AVX --half\_neigh 1

- INSTR\_RETIRED\_ANY
- FP\_ARITH\_INST\_RETIRED\_SCALAR\_DOUBLE.....(1)
- FP\_ARITH\_INST\_RETIRED\_128B\_PACKED\_DOUBLE.....(2)
- FP\_ARITH\_INST\_RETIRED\_256B\_PACKED\_DOUBLE.....(3)
- FP\_ARITH\_INST\_RETIRED\_512B\_PACKED\_DOUBLE.....(4)
- CPI (Cycle Per Instruction)
- Vectorization Ratio = [ (2) + (3) + (4) ] / [ (1) + (2) + (3) + (4) ]

# Whole Application Measurements

|  | HN Novec    | HN SSE      | HN AVX      | HN AVX512   | FN Novec    | FN SSE      | FN AVX      | FN AVX512   |
|--|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
| <b>Total Instructions</b>                    | 1.04419E+11 | 1.25319E+11 | 1.10110E+11 | 4.16818E+10 | 1.68247E+11 | 1.70410E+11 | 9.70627E+10 | 6.55391E+10 |
| <b>Arithmetic Instructions</b>               | 4.27175E+10 | 2.76718E+10 | 2.83674E+10 | 1.55086E+10 | 7.47834E+10 | 4.82002E+10 | 3.29768E+10 | 2.64749E+10 |
| <b>Arithmetic Percentage [%]</b>             | 40.91       | 22.08       | 25.76       | 37.21       | 44.45       | 28.28       | 33.97       | 40.40       |
| <b>Runtime (RDTSC) [s]</b>                   | 19.64       | 20.93       | 20.30       | 11.47       | 33.44       | 23.67       | 15.84       | 13.07       |
| <b>CPI</b>                                   | 0.5787      | 0.5157      | 0.5665      | 0.7620      | 0.6474      | 0.4358      | 0.4871      | 0.5708      |
| <b>Clock Frequency [MHz]</b>                 | 3235.4168   | 3256.9114   | 3238.7995   | 3018.8747   | 3354.7131   | 3285.0358   | 3177.8148   | 3079.8612   |
| <b>Vectorization Ratio [%]</b>               | 0           | 71.70       | 72.39       | 57.46       | 0           | 73.30       | 58.62       | 52.73       |
| <b>Instructions vs. HN [%]</b>               | -           | -           | -           | -           | 161.13      | 135.98      | 88.15       | 157.24      |
| <b>Instructions vs. Novec [%]</b>            | -           | 120.02      | 105.45      | 39.92       | -           | 101.29      | 57.69       | 38.95       |
| <b>Arithmetic Instructions vs. Novec [%]</b> | -           | 64.78       | 66.41       | 36.31       | -           | 64.45       | 44.10       | 35.40       |

| HN MiniMD: Function Name | Time(%) | Self Seconds | Calls |
|--------------------------|---------|--------------|-------|
| ForceLJ::compute         | 77.47   | 14.40        | 102   |
| Neighbor::build          | 17.21   | 3.20         | 6     |
| Integrate::run           | 1.72    | 0.32         | 1     |
| Atom::pack_comm          | 0.48    | 0.09         | 570   |
| Atom::unpack_reverse     | 0.38    | 0.07         | 612   |

| FN MiniMD: Function Name | Time(%) | Self Seconds | Calls |
|--------------------------|---------|--------------|-------|
| ForceLJ::compute         | 77.43   | 25.28        | 102   |
| Neighbor::build          | 19.57   | 6.39         | 6     |
| Integrate::run           | 1.26    | 0.41         | 1     |
| create_atoms             | 0.25    | 0.08         | 1     |
| Atom::pack_comm          | 0.18    | 0.06         | 570   |

## Input Parameters

- Sequential and no vectorization
- Lennard-Jones potential
- Cutoff distance: 2.5
- Neighbor skin width: 0.3
- Number of atoms:  $48 * 48 * 48 * 4 = 442368$
- Timesteps: 100
- Timestep size: 0.005



```
LIKWID_MARKER_START("HalfNeigh");
for(int i = 0; i < nlocal; i++) {
    neighs = &neighbor.neighbors[i * neighbor.maxneighs];
    const int numneighs = neighbor.numneigh[i];
    const MMD_float xtmp = x[i * PAD + 0];
    const MMD_float ytmp = x[i * PAD + 1];
    const MMD_float ztmp = x[i * PAD + 2];

    MMD_float fix = 0.0;
    MMD_float fiy = 0.0;
    MMD_float fiz = 0.0;

#ifdef USE_SIMD
#pragma simd reduction (+: fix,fiy,fiz)
#endif
    for(int k = 0; k < numneighs; k++) {
        const int j = neighs[k];
        const MMD_float delx = xtmp - x[j * PAD + 0];
        const MMD_float dely = ytmp - x[j * PAD + 1];
        const MMD_float delz = ztmp - x[j * PAD + 2];
        const MMD_float rsq = delx * delx + dely * dely + delz * delz;
```

```
    if(rsq < cutforcesq) {
        const MMD_float sr2 = 1.0 / rsq;
        const MMD_float sr6 = sr2 * sr2 * sr2 * sigma6;
        const MMD_float force = 48.0 * sr6 * (sr6 - 0.5) * sr2 * epsilon;

        fix += delx * force;
        fiy += dely * force;
        fiz += delz * force;

        if(j < nlocal) {
            f[j * PAD + 0] -= delx * force;
            f[j * PAD + 1] -= dely * force;
            f[j * PAD + 2] -= delz * force;
        }
    }
    f[i * PAD + 0] += fix;
    f[i * PAD + 1] += fiy;
    f[i * PAD + 2] += fiz;
}
LIKWID_MARKER_STOP("HalfNeigh");
```

```
LIKWID_MARKER_START("FullNeigh");
for(int i = 0; i < nlocal; i++) {
    neighs = &neighbor.neighbors[i * neighbor.maxneighs];
    const int numneighs = neighbor.numneigh[i];
    const MMD_float xtmp = x[i * PAD + 0];
    const MMD_float ytmp = x[i * PAD + 1];
    const MMD_float ztmp = x[i * PAD + 2];

    MMD_float fix = 0;
    MMD_float fiy = 0;
    MMD_float fiz = 0;

    for(int k = 0; k < numneighs; k++) {
        const int j = neighs[k];
        const MMD_float delx = xtmp - x[j * PAD + 0];
        const MMD_float dely = ytmp - x[j * PAD + 1];
        const MMD_float delz = ztmp - x[j * PAD + 2];
        const MMD_float rsq = delx * delx + dely * dely + delz * delz;
```

```
if(rsq < cutforcesq) {
    const MMD_float sr2 = 1.0 / rsq;
    const MMD_float sr6 = sr2 * sr2 * sr2 * sigma6;
    const MMD_float force = 48.0 * sr6 * (sr6 - 0.5) * sr2 * epsilon;

    fix += delx * force;
    fiy += dely * force;
    fiz += delz * force;
}

f[i * PAD + 0] += fix;
f[i * PAD + 1] += fiy;
f[i * PAD + 2] += fiz;
}
LIKWID_MARKER_STOP("FullNeigh");
```

# Hot Spots Measurements

|  | HN Novec    | HN SSE      | HN AVX      | HN AVX512   | FN Novec    | FN SSE      | FN AVX      | FN AVX512   |
|--|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
| <b>Total Instructions</b>                    | 7.59497E+10 | 9.70560E+10 | 8.58887E+10 | 1.75340E+10 | 1.16419E+11 | 1.18790E+11 | 5.31449E+10 | 2.16958E+10 |
| <b>Arithmetic Instructions</b>               | 3.59216E+10 | 2.10575E+10 | 2.18203E+10 | 8.99739E+9  | 6.21597E+10 | 3.57582E+10 | 2.06019E+10 | 1.41359E+10 |
| <b>Arithmetic Percentage [%]</b>             | 47.30       | 21.70       | 25.41       | 51.31       | 53.39       | 30.10       | 38.77       | 65.16       |
| <b>Runtime (RDTSC) [s]</b>                   | 14.71       | 15.78       | 16.12       | 6.32        | 25.10       | 15.39       | 7.82        | 4.88        |
| <b>CPI</b>                                   | 0.6242      | 0.5270      | 0.6083      | 1.0790      | 0.7230      | 0.4263      | 0.4710      | 0.6952      |
| <b>Clock Frequency [MHz]</b>                 | 3234.4792   | 3252.9333   | 3252.1095   | 3006.7474   | 3362.6897   | 3300.8917   | 3213.3516   | 3104.8627   |
| <b>Vectorization Ratio [%]</b>               | 0           | 93.50       | 93.72       | 98.49       | 0           | 98.38       | 93.41       | 98.42       |
| <b>Instructions vs. HN [%]</b>               | -           | -           | -           | -           | 153.28      | 122.39      | 61.88       | 123.74      |
| <b>Instructions vs. Novec [%]</b>            | -           | 127.79      | 113.09      | 23.09       | -           | 102.04      | 45.65       | 18.64       |
| <b>Arithmetic Instructions vs. Novec [%]</b> | -           | 58.62       | 60.74       | 25.05       | -           | 57.53       | 33.14       | 22.74       |

# Hot Spots Measurements

---

## Analysis of SIMD Vectorization Effectiveness

- When no SIMD vectorization is applied, the numbers of arithmetic instructions and total instructions of FN MiniMD are both larger than those numbers of HN MiniMD.
- For both HN and FN MiniMD, whichever SIMD vectorization is applied, the number of arithmetic instructions will always reduce. Especially for AVX512, that will reduce by about 75%.
- When no SIMD vectorization is applied, FN MiniMD costs more time than HN MiniMD. However, SIMD vectorization is much more effective for FN MiniMD, so that the runtime of FN MiniMD will reduce by about 80% at most with SIMD vectorization, while the runtime of HN MiniMD will only reduce by about 50% at most with SIMD vectorization.

---

# Thank You!

